1 Mutual Information

1.1 Exponential Decay in Markov Chains

If we have a Markov chain defined by the matrix $\boldsymbol{M}$, which is *irreducible* and *aperiodic*, and has a finite state space $E = \{1, \dots, n\}$, then we have that

$$\lim_{i \to \infty} \boldsymbol{M}^i = \boldsymbol{M}_{\boldsymbol{\mu}} \quad ,$$

where $\boldsymbol{M}_{\boldsymbol{\mu}}$ is the matrix whose columns all consist of the unique stationary probability distribution $\boldsymbol{\mu}$.

Now, let us consider two random variables X and Y, which will denote the state of the Markov chain at times t_0 and $t_0 + \tau$ respectively. We assume that we measure these variables very late in the process, where we already have that $\boldsymbol{M}^{t_0} \approx \boldsymbol{M}_{\boldsymbol{\mu}}$. We will use this "equality" later.

Our goal now is to quantify the mutual information of X and Y, that is, the discrepancy between the joint probability distribution $P(X, Y)$ and the one defined by the product of the two marginalized distributions, that is $P'(X, Y) \coloneqq P(X) \cdot P(Y)$. We use the Kullback-Leibler divergence, so our target expression becomes

$$D(P(X, Y) \parallel P'(X, Y)) \quad .$$

Note that of course this divergence $I(X, Y) \coloneqq D(P(X, Y) \parallel P'(X, Y))$ depends on the properties of M, as well as on τ. Because $\boldsymbol{M}$ is irreducible and aperiodic, it follows that $|\lambda_2| < 1$. The claim is that

$$I(X, Y) \in \mathcal{O}(|\lambda_2|^\tau) \quad .$$

There is a lot of math involved, so let us first get an intuition for what is going on. When considering Markov chains, we consider a set of states, say $E = \{A, B, C\}$, and for each time $t \in \mathbb{N}$ we assign a probability to the random variable $X_t \in E$. So let us consider the following Markov chain in figure 1.

If $\tau = 1$, i.e. we consider the mutual information of to consecutive states, we get a large value of $I(X, Y)$, as if X_{t_0} is either A or C, then X_{t_0+1} is uniquely determined, so we have a strong dependency between the two random variables. If, however, we have $\tau = 5$, then we can reach every state independent of the starting position. To see this, note that we can reach every state from A in four steps:

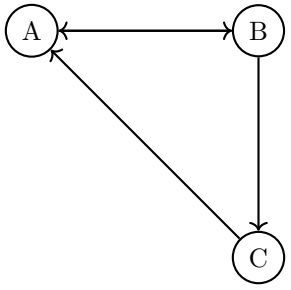

Figure 1: A simple irreducible aperiodic Markov chain. Note that if $X_{t_0} = C$, then we know that $X_{t_0+1} = A$.

- $A \to B \to C \to A \to B$
- $A \to B \to A \to B \to C$
- $A \to B \to A \to B \to A$

The last step can then be used to go around in a cycle. If we on the other hand started at B or C, then we could go to A in one step, and consequently to every other state in the following four. Hence, the probability distribution will "wash out" over time and converge to the stationary one, which results in a decline of $I(X, Y)$ for increasing τ.

Because we measure our X very late in time, meaning t_0 is very large, we will have that $P(X = a) \approx \boldsymbol{\mu}_a$ because of this "washing out". Similarly, we have $P(Y = b) \approx \boldsymbol{\mu}_b$, since the probability distribution will only get attracted more towards $\boldsymbol{\mu}$. As we now increase τ, $P(Y = b \,|\, X = a)$ itself will converge to $\boldsymbol{\mu}_b$ exactly due to the same "washing out" reason. Note that $P(Y = b \,|\, X = a) = (\boldsymbol{M}^\tau)_{b,a} \xrightarrow{\tau \to \infty} \boldsymbol{\mu}_b$. And, of course, if $P(X = a, Y = b) = P(X = a) \cdot P(Y = b \,|\, X = a) = \boldsymbol{\mu}_a \cdot \boldsymbol{\mu}_b$, we have $I(X, Y) = 0$. Hence, in a sense the theorem describes how fast $\boldsymbol{M}^\tau p_0$ converges to $\boldsymbol{\mu}$, or, equivalently, $\boldsymbol{M}^\tau$ towards $\boldsymbol{M}_{\boldsymbol{\mu}}$.

Now it's time to dive into the math. In the following, we try to reconstruct the arguments given in the paper. We also adopt the notation $P(a, b) \equiv P(X = a, Y = b)$. By definition of the Kullback-Leibler divergence, we have

$$D(P(X, Y) \,\|\, P'(X, Y)) = \sum_{(a,b) \in E^2} P(a, b) \log_B \frac{P(a, b)}{P(a)P(b)} \quad .$$

The idea is now that $log_B(\bullet)$ is *concave*. Hence, we can upper bound it by its

Taylor expansion of the first degree at the point $x_0 = 1$:

$$\begin{aligned}\log_B(x) &\leq \log_B(x_0) + \log_B'(x_0)(x - x_0) \\ &= 0 + \frac{\ln'(x_0)}{\ln(B)}(x-1) \\ &= \frac{\frac{1}{x_0}}{\ln(B)}(x-1) \\ &= \frac{x-1}{\ln(B)} \quad .\end{aligned}$$

For simplicity, we set $B \coloneqq e$. So our expression becomes

$$\begin{aligned}D(P(X,Y) \,\|\, P'(X,Y)) &\leq \frac{1}{\ln(B)} \sum_{(a,b)\in E^2} P(a,b)\left(\frac{P(a,b)}{P(a)P(b)} - 1\right) \\ &= \sum_{(a,b)\in E^2} P(a,b)\left(\frac{P(a,b)}{P(a)P(b)} - 1\right) \\ &= \left(\sum_{(a,b)\in E^2} P(a,b)\frac{P(a,b)}{P(a)P(b)}\right) - 1 \\ &= \left(\sum_{(a,b)\in E^2} \frac{P(a,b)^2}{P(a)P(b)}\right) - 1 \\ &=: I_R(X,Y) \quad .\end{aligned}$$

The authors of the paper coin this definition for $I_R(X,Y)$ the *rational mutual information*, as it has some useful properties. As discussed, we can approximate $P(a) \approx \boldsymbol{\mu}_a$ and $P(b) \approx \boldsymbol{\mu}_b$, and also $P(b \,|\, a) = (\boldsymbol{M}^\tau)_{b,a}$. Thus:

$$\begin{aligned}I_R(X,Y) + 1 &= \sum_{(a,b)\in E^2} \frac{P(a,b)^2}{P(a)P(b)} \\ &= \sum_{(a,b)\in E^2} \frac{P(b \,|\, a)^2 P(a)^2}{P(a)P(b)} \\ &= \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b}\left[(\boldsymbol{M}^\tau)_{b,a}\right]^2 \quad .\end{aligned}$$

Let us now focus on $(\boldsymbol{M}^\tau)_{b,a}$. For simplicity, we consider the case that the eigenvalues of $\boldsymbol{M}$ are all distinct, and hence $\boldsymbol{M}$ being diagonalizable. Note that since M is irreducible and aperiodic, we have that $1 = \lambda_1 > |\lambda_2| \geq \cdots \geq |\lambda_n|$. The authors provide proof for the other case as well. But for now, let

$$\boldsymbol{M} = \boldsymbol{B}\boldsymbol{D}\boldsymbol{B}^{-1}$$

be the diagonalization of $\boldsymbol{M}$. Of course, we immediately see that $\boldsymbol{M}^\tau = \boldsymbol{B}\boldsymbol{D}^\tau\boldsymbol{B}^{-1}$. Hence, it is easy to verify that

$$(\boldsymbol{M}^\tau)_{b,a} = \sum_{c=1}^{n} \lambda_c^\tau \boldsymbol{B}_{b,c}(\boldsymbol{B}^{-1})_{c,a} \quad .$$

Okay, that was a lot of math. Now it is a good time to reassure ourselves what we actually have achieved. What do we expect $(\boldsymbol{M}^\tau)_{b,a}$ to look like for $\tau \to \infty$? $\boldsymbol{\mu}_b$ of course. What does $\boldsymbol{B}$ look like? Well, this is very hard to tell, it at least should have a scaled version of $\boldsymbol{\mu}$ in its first column. But we cannot really infer any information about $\boldsymbol{B}^{-1}$. But we know

$$\begin{aligned}
\boldsymbol{\mu}_b &= \lim_{\tau\to\infty} (\boldsymbol{M}^\tau)_{b,a} \\
&= \lim_{\tau\to\infty} \sum_{c=1}^{n} \lambda_c^\tau \boldsymbol{B}_{b,c}(\boldsymbol{B}^{-1})_{c,a} \\
&= \lambda_1 \boldsymbol{B}_{b,1}(\boldsymbol{B}^{-1})_{1,a} \quad .
\end{aligned}$$

So we know that

$$(\boldsymbol{M}^\tau)_{b,a} = \boldsymbol{\mu}_b \pm \mathcal{O}(|\lambda_2|^\tau) \quad .$$

Note that this is informal writing. It would be more precise to state that $|(\boldsymbol{M}^\tau)_{b,a} - \boldsymbol{\mu}_b| \in \mathcal{O}(|\lambda_2|^\tau)$.

This is looking promising, as this means that the discrepancy between $(\boldsymbol{M}^\tau)_{b,a}$ and $\boldsymbol{\mu}_b$ decays exponentially. The only thing left to do is translating this exponential decay to the mutual independence measure $I_R(X,Y)$. To this end, we plug our results back into our previous equation. Note that this step deviates from the procedure in the paper (own interpretation, informal!). Thus:

$$\begin{aligned}
I_R(X,Y) &= \left(\sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} \left[(\boldsymbol{M}^\tau)_{b,a}\right]^2 \right) - 1 \\
&= \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} \left[(\boldsymbol{M}^\tau)_{b,a}\right]^2 - \boldsymbol{\mu}_a\boldsymbol{\mu}_b \\
&= \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} \left[\boldsymbol{\mu}_b \pm \mathcal{O}(|\lambda_2|^\tau)\right]^2 - \boldsymbol{\mu}_a\boldsymbol{\mu}_b \\
&= \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} \left[\boldsymbol{\mu}_b^2 \pm \mathcal{O}(|\lambda_2|^\tau)\right] - \boldsymbol{\mu}_a\boldsymbol{\mu}_b \\
&= \pm \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} \mathcal{O}(|\lambda_2|^\tau) \quad ,
\end{aligned}$$

where we have used multiple facts about $\boldsymbol{\mu}$. For instance, $\sum_{a\in E}\boldsymbol{\mu}_a = 1$ and thus $\sum_{(a,b)\in E^2}\boldsymbol{\mu}_a\boldsymbol{\mu}_b = 1$, as well as $0 < \boldsymbol{\mu}_a < 1$ for all $a \in E$ (at least for $|E| > 1$). We now use the latter inequality again: We see that we can always bound $\frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b}$ from above, i.e. there exists $\alpha \in \mathbb{R}$ s.t. for all $(a,b) \in E^2$ we have $\frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b} < \alpha$. Hence:

$$
\begin{aligned}
&|I_R(X,Y)| \in \sum_{(a,b)\in E^2} \frac{\boldsymbol{\mu}_a}{\boldsymbol{\mu}_b}\mathcal{O}(|\lambda_2|^\tau) \\
\implies\ &|I_R(X,Y)| \in \sum_{(a,b)\in E^2} \alpha\mathcal{O}(|\lambda_2|^\tau) \\
\implies\ &|I_R(X,Y)| \in n^2\alpha\mathcal{O}(|\lambda_2|^\tau) \\
\implies\ &|I_R(X,Y)| \in \mathcal{O}(|\lambda_2|^\tau) \quad .
\end{aligned}
$$

Of course, $I_R(X,Y) \geq 0$, so really $I_R(X,Y) \in \mathcal{O}(|\lambda_2|^\tau)$. Since $0 \leq I(X,Y) \leq I_R(X,Y)$, we also have $I(X,Y) \in \mathcal{O}(|\lambda_2|^\tau)$.